

/ AMD Tech Day

AI#3 - Vitis AI DPU: Explore a Vivado based flow targeting the Kria KV260

Dimitrios Kolosov, Senior Field Application Engineer - FPGA Specialist

 AVNET[®] SILICA

8th November, 2023



/ Agenda


- Overviews
- AMD Vivado Steps
- AMD PetaLinux Steps
- KV260 Deployment
- Q&A



Overviews

K26 SOM Overview

Based on the Zynq® UltraScale+™ MPSoC Architecture



COMPUTE

Application Processor	64-bit Quad-Core Arm® Cortex®-A53
Real-Time Processor	32-bit Dual-Core Arm Cortex-R5F
Graphics Processor	Arm Mali™-400MP2
Programmable Logic	256K System Logic Cells
Deep Learning Processor	4K INT8 (upgradable to INT4)
Video Codec (H.264/H.265)	Up to 32 Streams (total resolution ≤ 4Kp60)
Memory	26.6Mb On-Chip SRAM
Security	IEC62443 Security w/HW Root-of-Trust

INTERFACES

Camera	11 x4 Full MIPI or sub-LVDS Interfaces 1 x4 SLVS-EC Interfaces
USB	4x USB 2.0 / 3.0
Multi-Media	DisplayPort, HDMI
Network	1Gb up to 40Gb Ethernet (w/GigE Vision)
Memory Interface	4GB 64-bit DDR4
Transceivers	4x 12.5Gb/s, 4x 6Gb/s
Mechanical	77 x 60 x 11mm w/ dual 240-pin connectors

1.4 TOPs

Comparing Kria™ K24 vs. K26 SOM



\$250 | \$350
(C-grade | I-grade)



CONNECTOR
COMPATIBLE



\$325 | \$450
(C-grade | I-grade)

Cost-optimized SOM for lower power, smaller
form-factor & cost sensitive industrial
applications

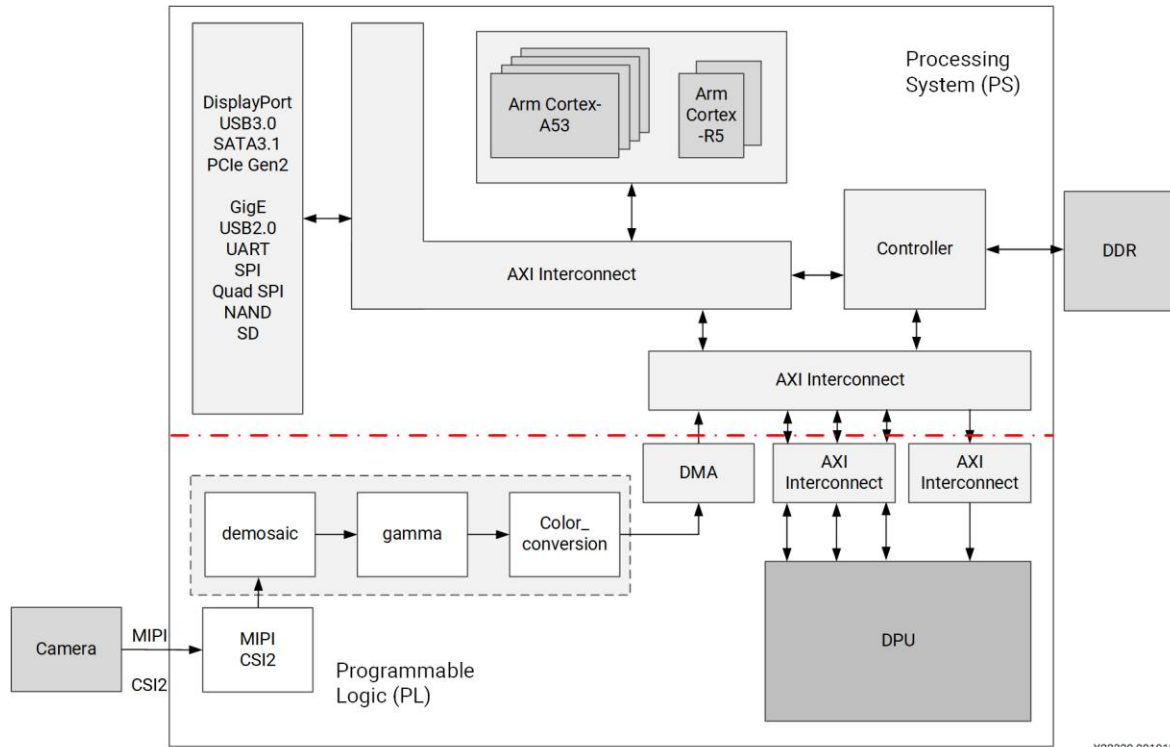
Mid-range SOM for Vision AI and Robotics
applications requiring higher performance per
watt

	K24 SOM		K26 SOM
SILICON (SYS LOGIC CELLS)	XCK24 InFO (154K)	▶ SILICON ▶	XCK26 (256K)
SOM I/O ACCESS	1x 240-Pin Connector, 1x 40-Pin Connector	▶ SOM I/O ▶	2x 240-Pin Connectors
FORM FACTOR	60 x 42mm	▶ 46% SMALL ▶	60 x 77mm
MEMORY	2GB LPDDR4 ¹ , 32 GB	▶ DDR, eMMC ▶	4GB DDR4, 16 GB
POWER ²	2.5W	▶ 51% LESS ▶	5.1W
STARTER KITS	KD240 DRIVES	▶ DEV KITS ▶	KV260 VISION AI, KR260 ROBOTICS

¹ ECC support available on K24 SOM I-grade

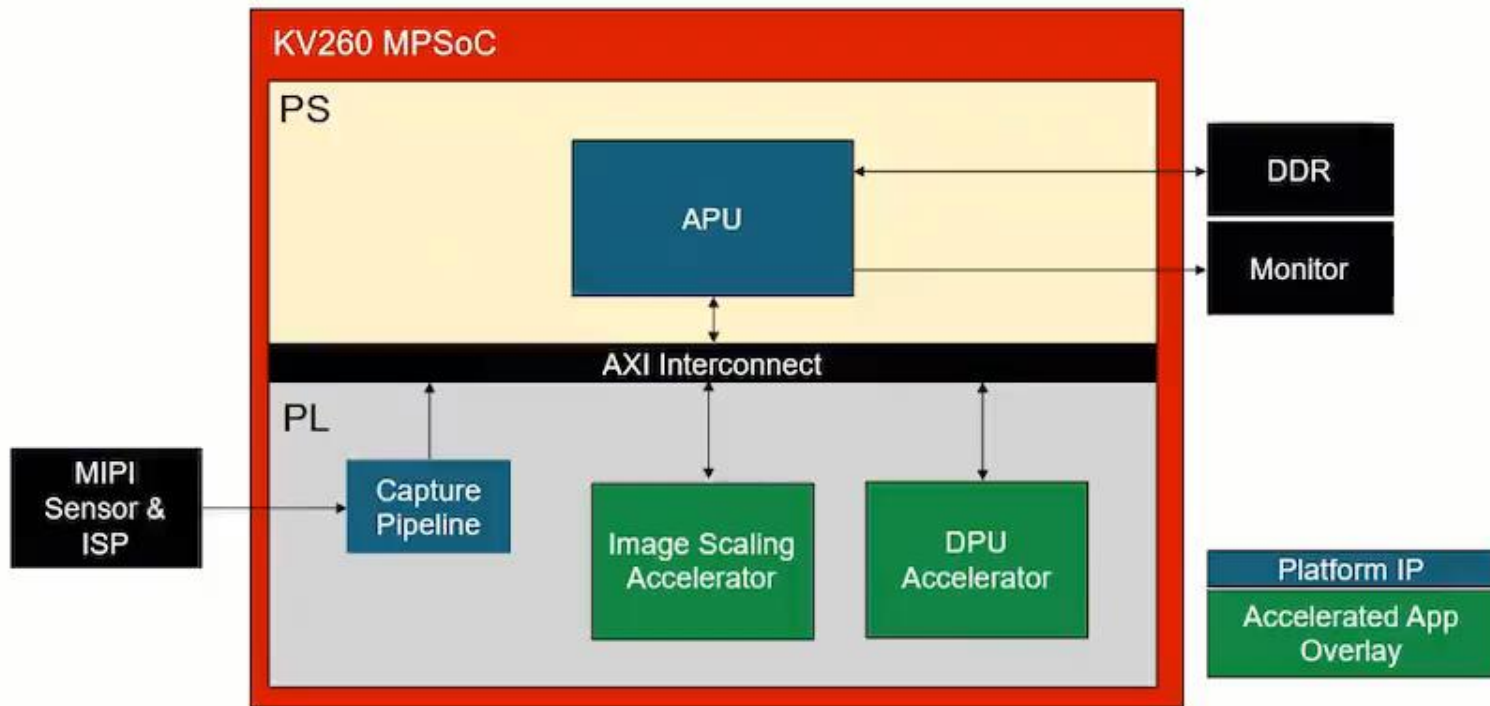
² Measured power while loading application specific bitstream on the SOM-based starter kit

/ Typical AI Example System



X22329-081919

/ Example of Dataflow



/ Zynq™ Ultrascale+™ DPU Compatibility Matrix

Vitis AI Release Version	DPUCZDX8G IP Version	Software Tools Version	Linux Kernel Version Tested
v3.5	4.1 (not updated*)	Vivado / Vitis / PetaLinux 2023.1	6.1
v3.0	4.1	Vivado / Vitis / PetaLinux 2022.2	5.15
v2.5	4.0	Vivado / Vitis / PetaLinux 2022.1	5.15
v2.0	3.4	Vivado / Vitis / PetaLinux 2021.2	5.10
v1.4	3.3	Vivado / Vitis / PetaLinux 2021.1	5.10
v1.3	3.3	Vivado / Vitis / PetaLinux 2020.2	5.4

Vitis AI Release Version	DPUCZDX8G IP Version	Software Tools Version	Linux Kernel Version Tested
v1.2	3.2	Vivado / Vitis / PetaLinux 2020.1	5.4
v1.1	3.2	Vivado / Vitis / PetaLinux 2019.2	4.19
v1.0	3.1	Vivado / Vitis / PetaLinux 2019.1	4.19
N/A (DNNDK)	3.0	Vivado / Vitis / PetaLinux 2019.1	4.19
N/A (DNNDK)	2.0	Vivado / Vitis / PetaLinux 2018.2	4.14
First Release (DNNDK)	1.0	Vivado / Vitis / PetaLinux 2018.1	4.14



AMD Vivado Steps

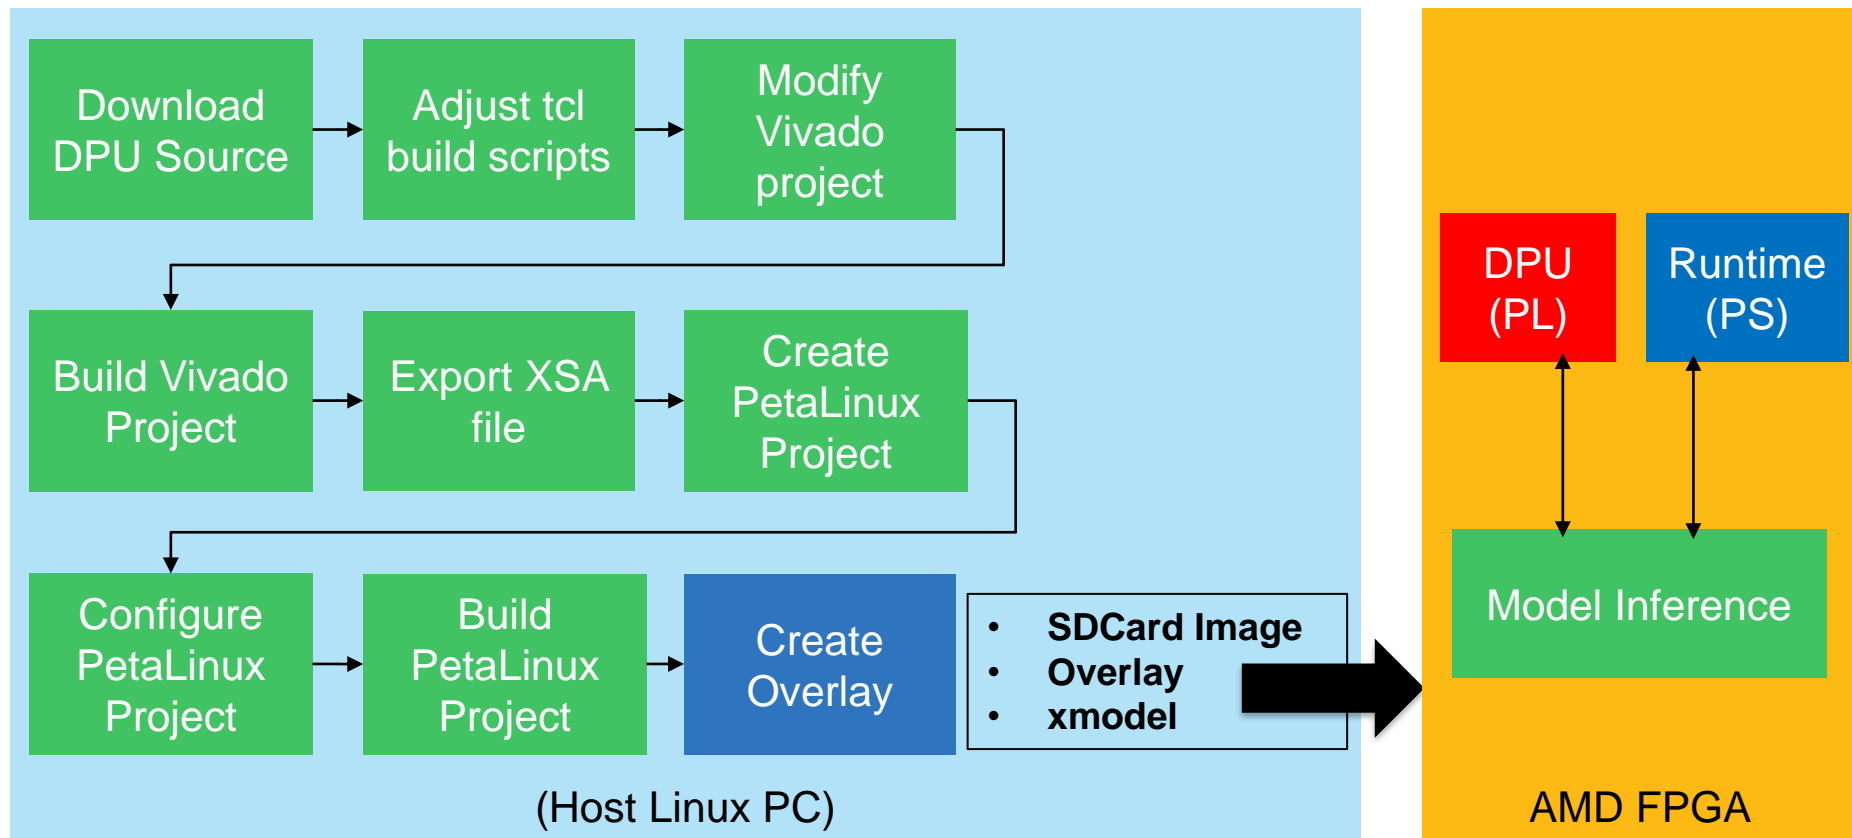


/ Goals

- Download, adjust and execute AMD-Xilinx tcl scripts for custom Vivado project with following DPU settings (vs pre-build petalinux with DPU images):
 - DPU Clock: 325MHz (300MHz originally)
 - DSP Clock: 650MHz (600MHz originally)
 - DPU Variant: B4096
 - DPU Cores: x1
 - UltraRAM: 50/64
- Create and configure PetaLinux project
- Create overlay
- Integrate and run custom overlay on KV260 Starter Kit

Original ref design
targets ZCU102
development board

/ AMD Vivado Flow (KV260)



/ Requirements:

- Vivado 2022.2 (Windows or Linux OS)
- Petalinux 2022.2 (Linux OS)
- KV260 Starter Kit BSP v2022.2 from PetaLinux Download page ([link](#))
- DPUCZDX8G_VAI_v3.0 (v4.1) source files from Vitis-AI GitHub ([link](#))

/ Download DPU IP

Vitis-AI/dpu at v3.0 · Xilinx/Vitis-AI · GitHub

Files

v3.0

Go to file

board_setup

demos

docker

docs

docsrc

dpu

ref_design_docs

README.md

version_compatibility.md

examples

model_zoo

src

third_party

CONTAINER-CPU-LICENSE

CONTAINER-ROCM-LICENSE

LICENSE

README.md

docker_run.sh

Vitis-AI / dpu /

documentation. For convenience, a separate IP repo is provided for users who do not wish to download the reference design. The IP is thus included both in the reference design, but also is available as a separate download.

Please refer to [DPU Nomenclature](#) for detailed information on the capabilities and device targets for Xilinx DPUs.

Edge IP

IP Name	Supported Platforms	Description	Reference Design	Product Guide	Read Me	IP-only Download
DPUCZDX8G	MPSoc / Kria K26	Programmable logic based DPU, targeting general purpose CNN inference with full support for the Vitis AI ModelZoo. Supports either the Vitis or Vivado flows on 16nm Zynq® UltraScale+™ platforms.	Download	PG338	Link	Get IP
DPUCVDX8G	VCK190	AIE centric DPU (requires some programmable logic), targeting general purpose CNN inference with full support for the Vitis AI ModelZoo. Supports the Vitis flow for 7nm Versal™ ACAPs.	Download	PG389	Link	Get IP
DPUCV2DX8G	VEK280	AIE-ML centric DPU (requires some programmable logic), targeting general purpose CNN inference with full support for the Vitis AI ModelZoo. Supports the Vitis flow for 7nm Versal AI Edge targets leveraging the AIE-ML architecture.	Early Access	---	---	---

AVNET SILICA

/ Adjust tcl build scripts

Edit ~/DPUCZDX8G_VAI_v3.0/prj/Vivado/hw/scripts/trd_prj.tcl

```
22 #*****
23 # set project
24 #*****
25 dict set dict_prj dict_sys prj_name      {AVNET KV260 DPU}
26 dict set dict_prj dict_sys prj_part     {xck26-sfvc784-2LV-c}
27 dict set dict_prj dict_sys prj_board    {KV260}
28
29 #*****
30 # set bd
31 # for bd_ooc: None for global, Hierarchical for ooc per IP
32 #*****
33 dict set dict_prj dict_sys bd_name      top
34 dict set dict_prj dict_sys bd_ooc      None
35
36 #*****
37 # set param
38 #*****
39 dict set dict_prj dict_param DPU_CLK_MHz {325}
40 dict set dict_prj dict_param REG_CLK_MHz {100}
41
42 #The following parameters correspond to Arch Tab of the IP GUI
43 dict set dict_prj dict_param DPU_NUM      {1}
44 dict set dict_prj dict_param DPU_ARCH     {4096}
45 dict set dict_prj dict_param DPU_RAM_USAGE {low}
46 dict set dict_prj dict_param DPU_CHN_AUG_ENA {1}
47 dict set dict_prj dict_param DPU_SAVE_ARGMAX_ENA {1}
48 dict set dict_prj dict_param DPU_CONV_RELU_TYPE {3}
49 dict set dict_prj dict_param DPU_ALU_PARALLEL_USER {4}
50 dict set dict_prj dict_param DPU_ALU_LEAKYRELU {0}
51 dict set dict_prj dict_param DPU_SFM_NUM  {0}
52
53 #The following parameters correspond to Advanced Tab of the IP GUI
54 dict set dict_prj dict_param DPU_SAXICLK_INDPD {1}
55 dict set dict_prj dict_param DPU_CLK_GATING_ENA {1}
56 dict set dict_prj dict_param DPU_DSP48_MAX_CASC_LEN {4}
57 dict set dict_prj dict_param DPU_DSP48_USAGE {high}
58 dict set dict_prj dict_param DPU_URAM_PER_DPU {50}
59
```

Project
settings

DPU
settings

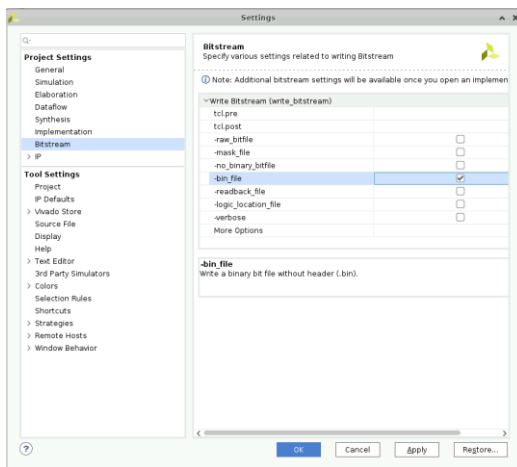
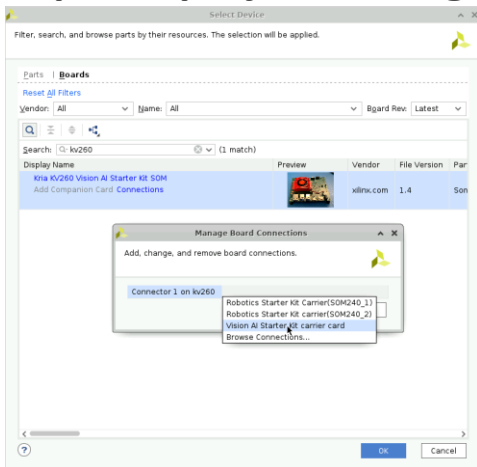
/ Modify Vivado settings

Create Vivado project via tcl script:

`/opt/Xilinx/Vivado/2022.2/bin/vivado -mode tcl -source trd_prj.tcl`

Vivado Project modifications:

- Update project settings

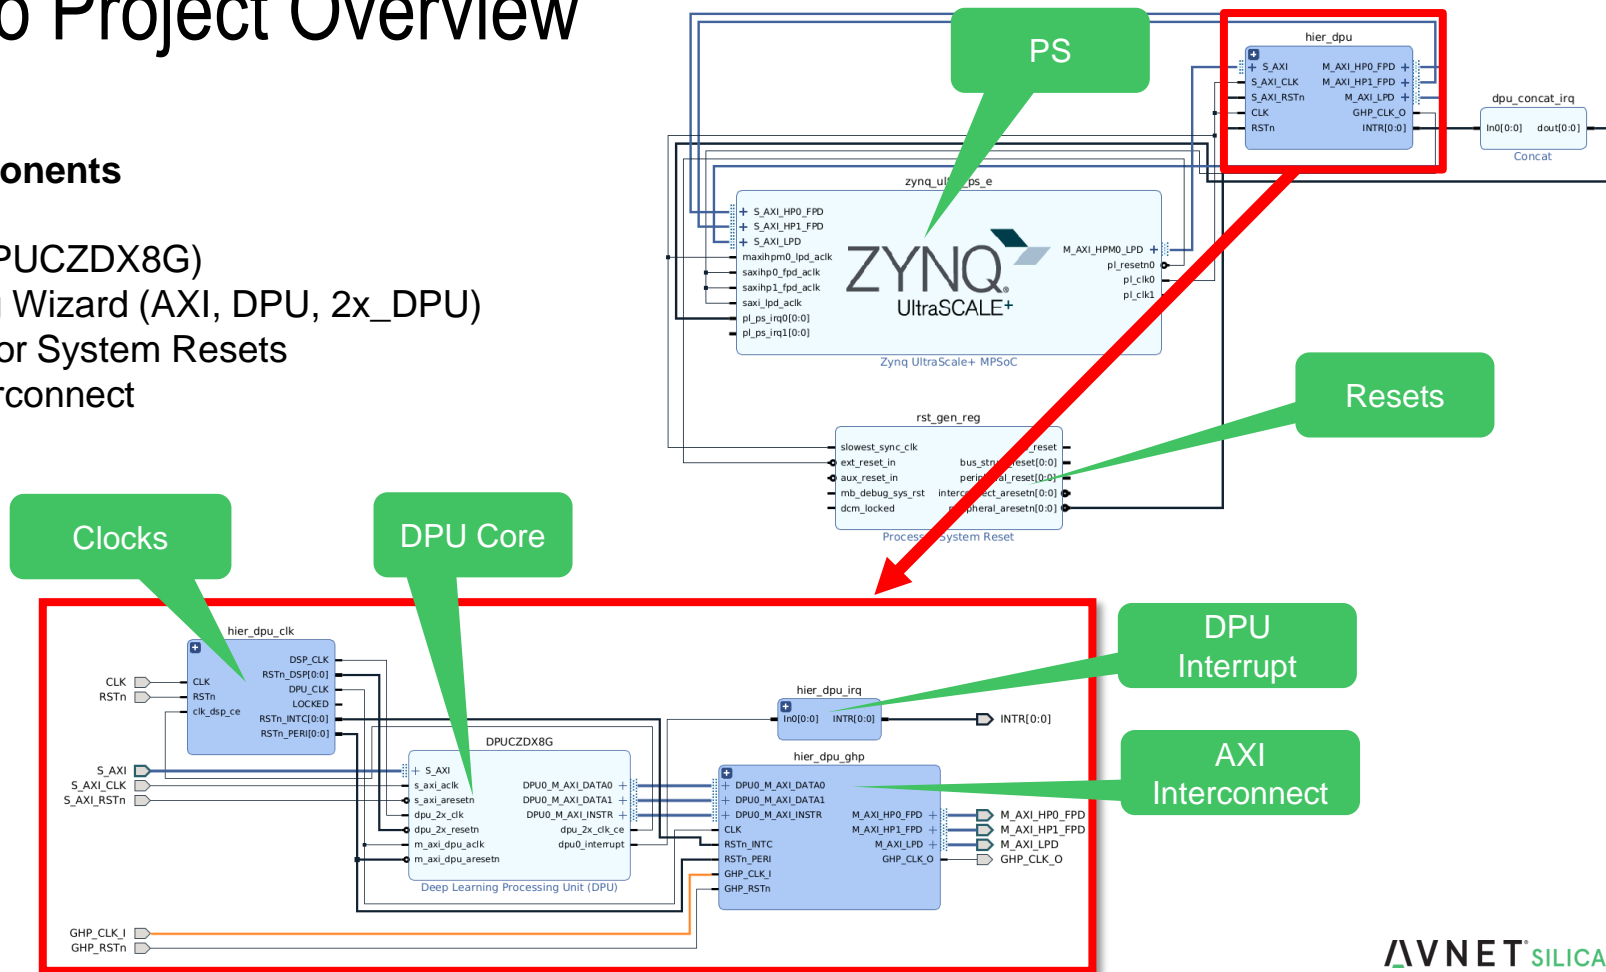


- Update block diagram IPs
- Update Implementation Strategy: Performance_Explore

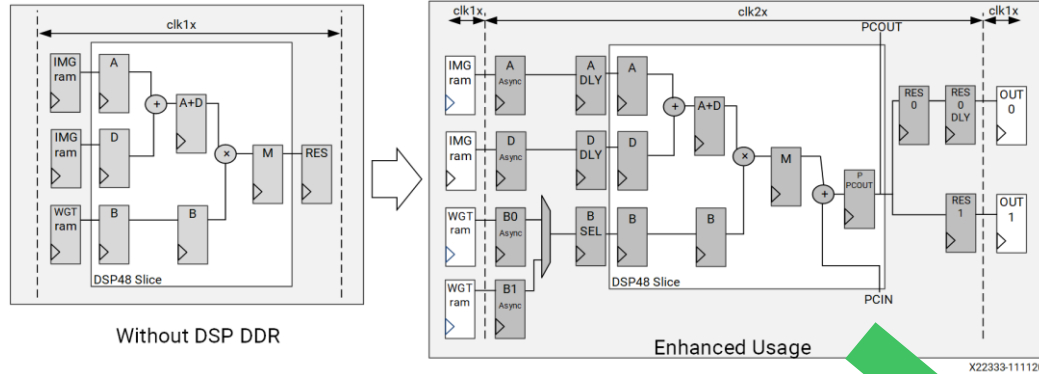
Vivado Project Overview

Core Components

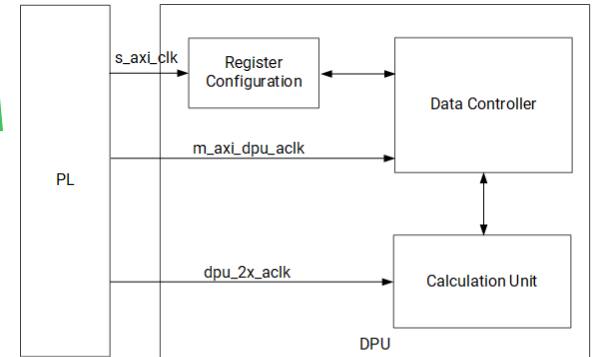
- PS
- DPU (DPUCZDX8G)
- Clocking Wizard (AXI, DPU, 2x_DPU)
- Processor System Resets
- AXI Interconnect



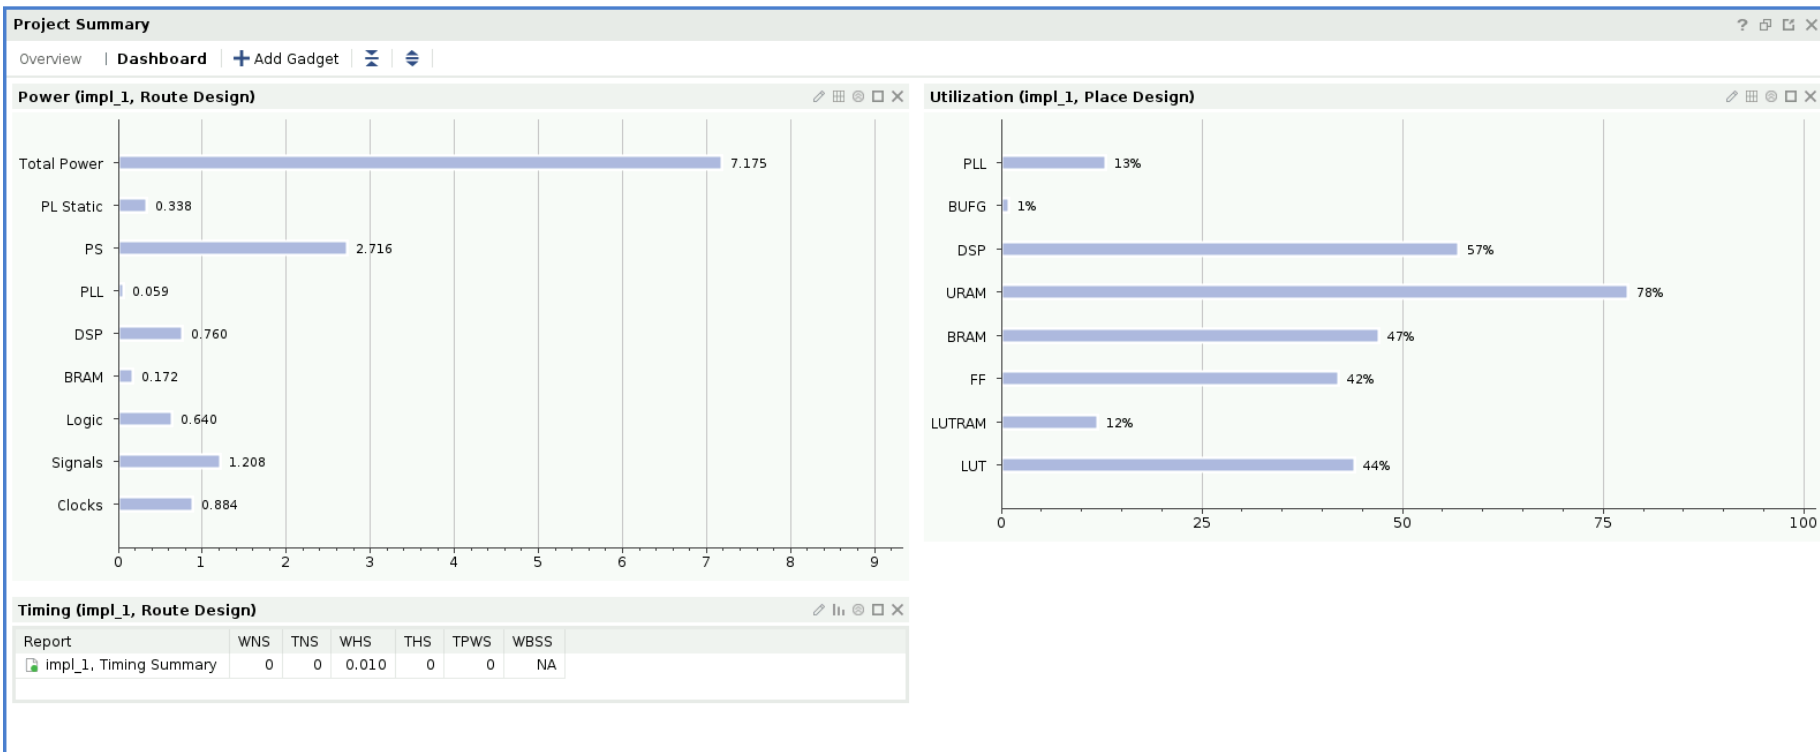
/ DPUCZDX8G Enhanced Usage



A DSP Double Data Rate (DDR) technique is used to boost the performance.



Vivado Project Compilation Results



/ Things to explore?

- Experimenting various DPU / clock frequency customizations

Table: Parallelism for Different Convolution Architectures

DPU Architecture	Pixel Parallelism (PP)	Input Channel Parallelism (ICP)	Output Channel Parallelism (OCP)	Peak Ops (operations/per clock)
B512	4	8	8	512
B800	4	10	10	800
B1024	8	8	8	1024
B1152	4	12	12	1150
B1600	8	10	10	1600
B2304	8	12	12	2304
B3136	8	14	14	3136
B4096	8	16	16	4096

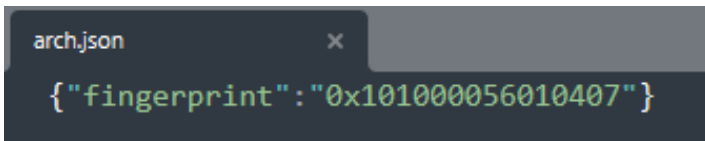
$$\text{GOPS} = \text{PP} * \text{ICP} * \text{OCP} * 2 * \text{Freq}$$

- Resource Utilizations
- Power Reports

/ Vivado required output files

1. arch.json file

Required to compile models in Vitis-AI Docker, targeting this DPU configuration, found in: `../DPUCZDX8G_VAI_v3.0/prj/Vivado/hw/srcs/top/ip/top_DPUCZDX8G_0/arch.json`

A screenshot of a code editor window with a dark background. The title bar shows 'arch.json' and a close button. The code content is a single line of JSON: `{"fingerprint": "0x101000056010407"}`.

```
arch.json
{"fingerprint": "0x101000056010407"}
```

2. *.xsa file

Hardware information of the Vivado design, to be used to create the PetaLinux project in following steps., found in: `../DPUCZDX8G_VAI_v3.0/prj/Vivado/hw/prj/*.xsa`

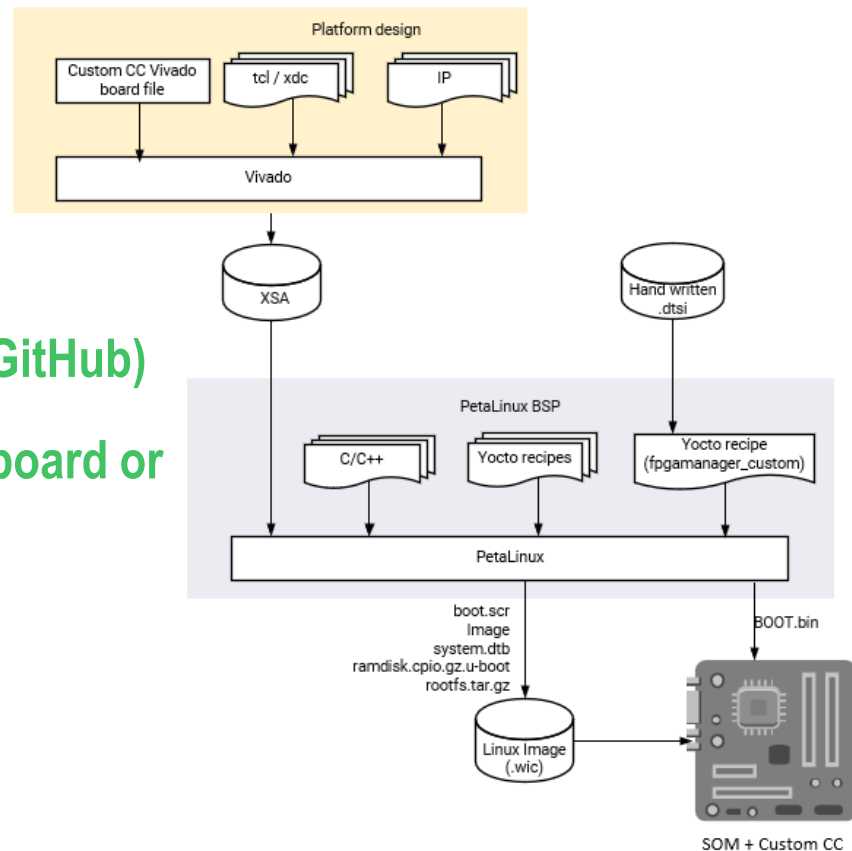


AMD PetaLinux Steps



/ PetaLinux Flow

1. v2022.2 includes PetaLinux build script
2. Previous versions, included steps (e.g. GitHub)
3. Script needs to be adjusted for custom board or in this case KV260 Starter Kit
(adjusted script “kv260_bsp_help.sh”
can be provided by Avnet Silica)



Summary of changes in next slides

/ PetaLinux Development Flow in short

Customize Linux in five steps:

- **petalinux-create**
 - Create a new project or components
- **petalinux-config**
 - Configure the project or specified components
 - Kernel, file system, libraries, debug options
 - Get hardware description from the Vivado® IDE
- **petalinux-build**
 - Build project or specified components
 - Kernel, file system, project-specific components
- **petalinux-boot**
 - Import FSBL and bitstream
 - Build image.ub, BOOT.BIN, etc.
 - Deploy to QEMU or hardware
- **petalinux-package**
 - Package custom image formats, firmware, prebuilt images & BSP for distribution



/ Create PetaLinux Project

1. Create PetaLinux Project based on BSP

```
petalinux-create -t project -s $bsp -n $prj_name
```

2. Update PetaLinux project with custom XSA

```
petalinux-config --get-hw-description=${xsa_dir} --silentconfig
```


/ Core PetaLinux Customizations for KV260

1. petalinux-config

Enable FPGA Manager

2. petalinux-config -c kernel

Enable DPU Driver

3. petalinux-config -c rootfs

Disable ZOCL and XRT (for use with Vitis Acceleration flow)

Enable auto-log in

Enable image feature package management

/ Create overlay for KV260

1. Device Tree (DT) Creation (in XSCT)

```
createdts -hw <XSA_FILE>.xsa -zocl -platform-name <PLATFORM NAME>  
-git-branch <xlnx_rel_vXXXX> -overlay -compile -out <OUTPUT_NAME>
```

2. Device Tree (DT) Compilation

```
dtc -@ -O dtb -o <OUTPUT>.dtbo <OUTPUT_NAME>/<OUTPUT_NAME>/<PLATFORM  
NAME>/psu_cortexa53_0/device_tree_domain/bsp/pl.dtsi
```

3. Create shell.json file:

```
echo '{ "shell_type" : "XRT_FLAT", "num_slots": "1" }' > shell.json
```

OVERLAY

- kv260.bit.bin (PL configuration file)
- kv260.dtbo (compiled DT)
- shell.json (overlay shell file)

/ Build PetaLinux Project

1. Build PetaLinux

```
petalinux-config -c kernel --silentconfig  
petalinux-config -c rootfs --silentconfig  
petalinux-build
```

2. Create SD Card

```
petalinux-package --wic --images-dir images/linux/ --bootfiles "ramdisk.cpio.gz.uboot, boot.scr,  
Image, system.dtb, system-zynqmp-sck-kv-g-revB.dtb" --disk-name "mmcblk1" --wic-extra-args "-c  
gzip"
```

KV260 Deployment

/KV260 – Load Overlay

- Copy overlay folder “KV260_DPU” to KV260 board

```
sudo mv ./KV260_DPU /lib/firmware/xilinx/
```

- Check available apps:

```
sudo xmutil listapps
```

```
root@avnet-kv260-dpu:~# sudo xmutil listapps
```

Accelerator	Accel_type	Base	Base_type	#slots(PL+AIE)	Active_slot
k26-starter-kits	XRT_FLAT	k26-starter-kits	XRT_FLAT	<0+0>	0
KV260_DPU	XRT_FLAT	KV260_DPU	XRT_FLAT	<0+0>	-1

```
root@avnet-kv260-dpu:~#
```

- Unload current app:

```
sudo xmutil unloadapp
```

```
root@avnet-kv260-dpu:~# sudo xmutil unloadapp
remove from slot 0 returns: 0 (Ok)
root@avnet-kv260-dpu:~# █
```

/ KV260 – Load Overlay

- Load our custom DPU app:
sudo xmutil loadapp KV260_DPU

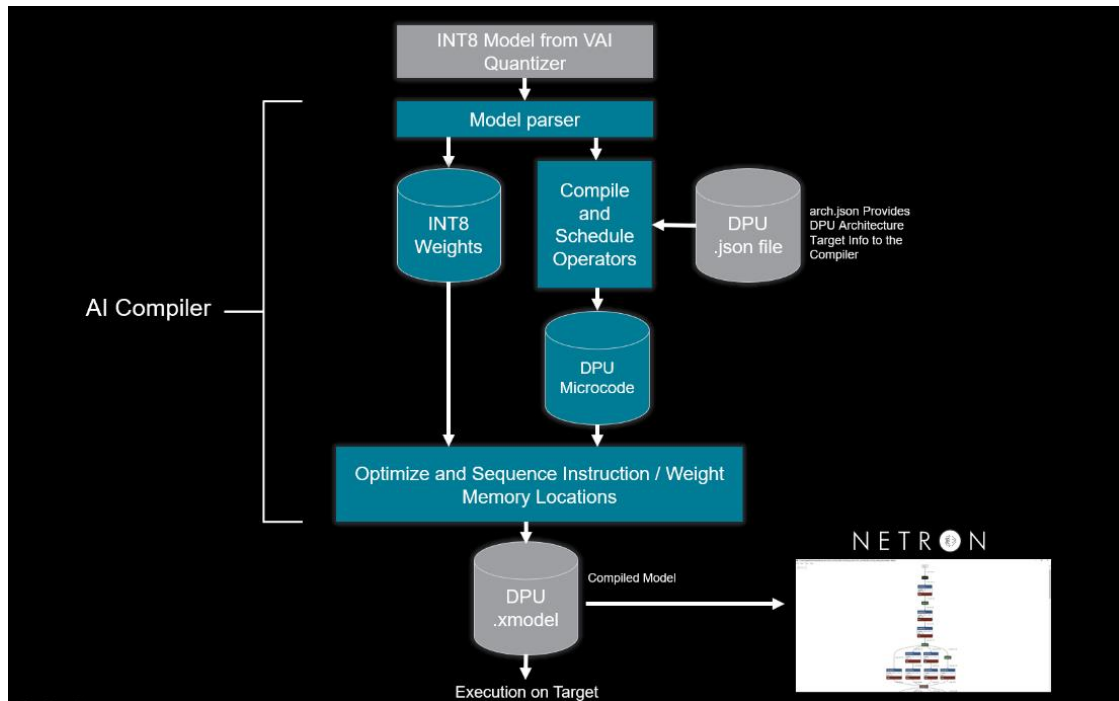
```
root@avnet-kv260-dpu:~#  
root@avnet-kv260-dpu:~# sudo xmutil loadapp KV260_DPU  
Nov 19 09:21:55 avnet-kv260-dpu kernel: OF: overlay: WARNING: memory leak will occur if overlay removed, property: /fpga-full/firmware-name  
Nov 19 09:21:55 avnet-kv260-dpu kernel: OF: overlay: WARNING: memory leak will occur if overlay removed, property: /fpga-full/pid  
Nov 19 09:21:55 avnet-kv260-dpu kernel: OF: overlay: WARNING: memory leak will occur if overlay removed, property: /fpga-full/resets  
Nov 19 09:21:55 avnet-kv260-dpu kernel: OF: overlay: WARNING: memory leak will occur if overlay removed, property: /fpga-full/uid  
Nov 19 09:21:55 avnet-kv260-dpu kernel: OF: overlay: WARNING: memory leak will occur if overlay removed, property: /__symbols__/overlay0  
Nov 19 09:21:55 avnet-kv260-dpu kernel: OF: overlay: WARNING: memory leak will occur if overlay removed, property: /__symbols__/overlay1  
Nov 19 09:21:55 avnet-kv260-dpu kernel: OF: overlay: WARNING: memory leak will occur if overlay removed, property: /__symbols__/afi0  
Nov 19 09:21:55 avnet-kv260-dpu kernel: OF: overlay: WARNING: memory leak will occur if overlay removed, property: /__symbols__/clocking0  
Nov 19 09:21:55 avnet-kv260-dpu kernel: OF: overlay: WARNING: memory leak will occur if overlay removed, property: /__symbols__/clocking1  
Nov 19 09:21:55 avnet-kv260-dpu kernel: OF: overlay: WARNING: memory leak will occur if overlay removed, property: /__symbols__/overlay2  
Nov 19 09:21:55 avnet-kv260-dpu kernel: OF: overlay: WARNING: memory leak will occur if overlay removed, property: /__symbols__/hier_dpu_DPUCZD8G  
Nov 19 09:21:55 avnet-kv260-dpu kernel: OF: overlay: WARNING: memory leak will occur if overlay removed, property: /__symbols__/misc_clk_0  
Nov 19 09:21:55 avnet-kv260-dpu kernel: OF: overlay: WARNING: memory leak will occur if overlay removed, property: /__symbols__/misc_clk_1  
KV260_DPU: loaded to slot 0  
root@avnet-kv260-dpu:~# █
```

/ Compile for target AMD Hardware

```
root@avnet-kv260-dpu:~# show_dpu  
device_core_id=0 device= 0 core = 0 fingerprint = 0x101000056010407 batch = 1 full_cu_name=unknown:dpu0
```

Requirements:

- Quantized Model
- DPU *.json file



/ KV260 – Quick DPU Check

Check DPU settings with: `xdputil query`

IP Version

cores

VAI Version

DPU
Architecture

DPU
Frequency

Fingerprint
ID

```
root@avnet-kv260-dpu:~# xdputil query
{
  "DPU IP Spec": {
    "DPU Core Count": 1,
    "IP version": "v4.1.0",
    "enable softmax": "False"
  },
  "VAI Version": {
    "libvart-runner.so": "Xilinx vart-runner Version: 3.0.0-c5d2bd43d951c174185d728b8e5bcda3869e0b39 2023-01-27-17:53:09 ",
    "libvitis_ai_library-dpu_task.so": "Xilinx vitis_ai_library dpu_task Version: 3.0.0-c5d2bd43d951c174185d728b8e5bcda3869e0b39 2023-01-13 06:58:30 [UTC]",
    "libxir.so": "Xilinx xir Version: xir-c5d2bd43d951c174185d728b8e5bcda3869e0b39 2023-01-27-17:52:29",
    "target_factory": "target-factory.3.0.0 c5d2bd43d951c174185d728b8e5bcda3869e0b39"
  },
  "kernels": [
    {
      "DPU Arch": "DPUCZDX8G_ISA1_B4096",
      "DPU Frequency (MHz)": 325,
      "XRT Frequency (MHz)": 100,
      "cu_idx": 0,
      "fingerprint": "0x101000056010407",
      "is_ovado_flow": true,
      "name": "DPU Core 0"
    }
  ]
}
```


/ KV260 – Quick DPU Benchmark

```
root@avnet-kv260-dpu:~# xdputil benchmark resnet50.xmodel 1
WARNING: Logging before InitGoogleLogging() is written to STDERR
I1119 09:22:41.180843 1625 test_dpu_runner_mt.cpp:477] shuffle results for batch...
I1119 09:22:41.181735 1625 performance_test.hpp:73] 0% ...
I1119 09:22:47.181915 1625 performance_test.hpp:76] 10% ...
I1119 09:22:53.182128 1625 performance_test.hpp:76] 20% ...
I1119 09:22:59.182329 1625 performance_test.hpp:76] 30% ...
I1119 09:23:05.182534 1625 performance_test.hpp:76] 40% ...
I1119 09:23:11.182749 1625 performance_test.hpp:76] 50% ...
I1119 09:23:17.182960 1625 performance_test.hpp:76] 60% ...
I1119 09:23:23.183177 1625 performance_test.hpp:76] 70% ...
I1119 09:23:29.183406 1625 performance_test.hpp:76] 80% ...
I1119 09:23:35.183670 1625 performance_test.hpp:76] 90% ...
I1119 09:23:41.183892 1625 performance_test.hpp:76] 100% ...
I1119 09:23:41.184017 1625 performance_test.hpp:79] stop and waiting for all threads terminated....
I1119 09:23:41.192155 1625 performance_test.hpp:85] thread-0 processes 6310 frames
I1119 09:23:41.192206 1625 performance_test.hpp:93] it takes 6155 us for shutdown
I1119 09:23:41.192234 1625 performance_test.hpp:94] FPS= 105.148 number of frames= 6310 time= 60.0105 seconds.
I1119 09:23:41.192289 1625 performance_test.hpp:96] BYEBYE
Test PASS.
```

+6% FPS by increasing the clock rate from 300MHz to 325MHz for ResNet50

DPU Clock / DSP Clock (MHz)	FPS
300 / 600	98.9
325 / 650	105.1



Q&A



Typical Questions

/ Are all the components of Vitis AI free?

Yes!

As of the 3.5 release all components are free!

For releases <3.5, the Vitis AI Optimizer does require a separate license which can be obtained free-of-charge upon request.

/ Is Vitis AI a separate download?

Yes! Users can get started by cloning the Vitis AI GitHub repository:

GitHub - Xilinx/Vitis-AI: Vitis AI is Xilinx's development stack for AI inference on Xilinx hardware platforms, including both edge devices and Alveo cards.

/ What Xilinx Target Device Families and Platforms does Vitis AI Support?

Vitis AI DPUs are available for:

- Zynq 7000
- Zynq Ultrascale+ MPSoC
- Versal AI Edge
- Versal AI Core

/ How does FPGA compare to CPU and GPU acceleration?

FPGA accelerated networks can run up to 90x faster as compared to CPU. FPGA accelerated networks are on par with GPU accelerated networks for throughput critical applications yet provide support for more custom applications.

FPGA accelerated networks are far superior to GPU accelerated networks for latency critical applications such as autonomous driving.

Lot of scope in reducing on a system level power / cost and creating scalable designs.

/ Is it possible to deploy the DPUCZ using Yocto flows, or even Ubuntu, rather than PetaLinux?

Yes!

What is important to consider is that each release of the Vitis AI tool and the DPUCZ IP is provided with drivers and a runtime that targets a specific Linux kernel release. Misalignment between the target kernel version can pose challenges and may require extensive code changes.

/ My DPU implementation does not meet my latency/throughput targets. Is there anything else I can do?

Yes!

Besides modifying architecture and/or taking advantage of pruning within Vitis AI, you can also explore FINN.

FINN implements each layer of a neural network separately and creates like this a custom very low latency and high throughput design in the PL.

The background is a solid green color. There are two white diagonal lines: one in the top-left corner and one in the bottom-right corner, both extending from the edge towards the center.

Thank You